

# **7a. Instruction Set Architecture –**

- translation software**
- floating point representation**

---

**EECS 370 – Introduction to Computer Organization – Winter 2018**

**Mark Brehob, Reetu Das, Harry Davis**

**EECS Department  
University of Michigan in Ann Arbor, USA**

**© Davis, Das, and Brehob, 2018**

The material in this presentation cannot be  
copied in any form without our written permission

# Announcements

---

- ❑ HW 2 is due Tuesday January 30<sup>th</sup> by 11:55 pm
- ❑ Project 1m and 1s are due Thursday February 1<sup>st</sup>.
- ❑ Midterm Exam
  - 8:15pm (sharp) to 10:15pm on February 22<sup>nd</sup>
  - You have until January 30<sup>th</sup> to let us know of any conflict you may have.
    - Directions for doing so are on the website under “administrative requests”

# Review

---

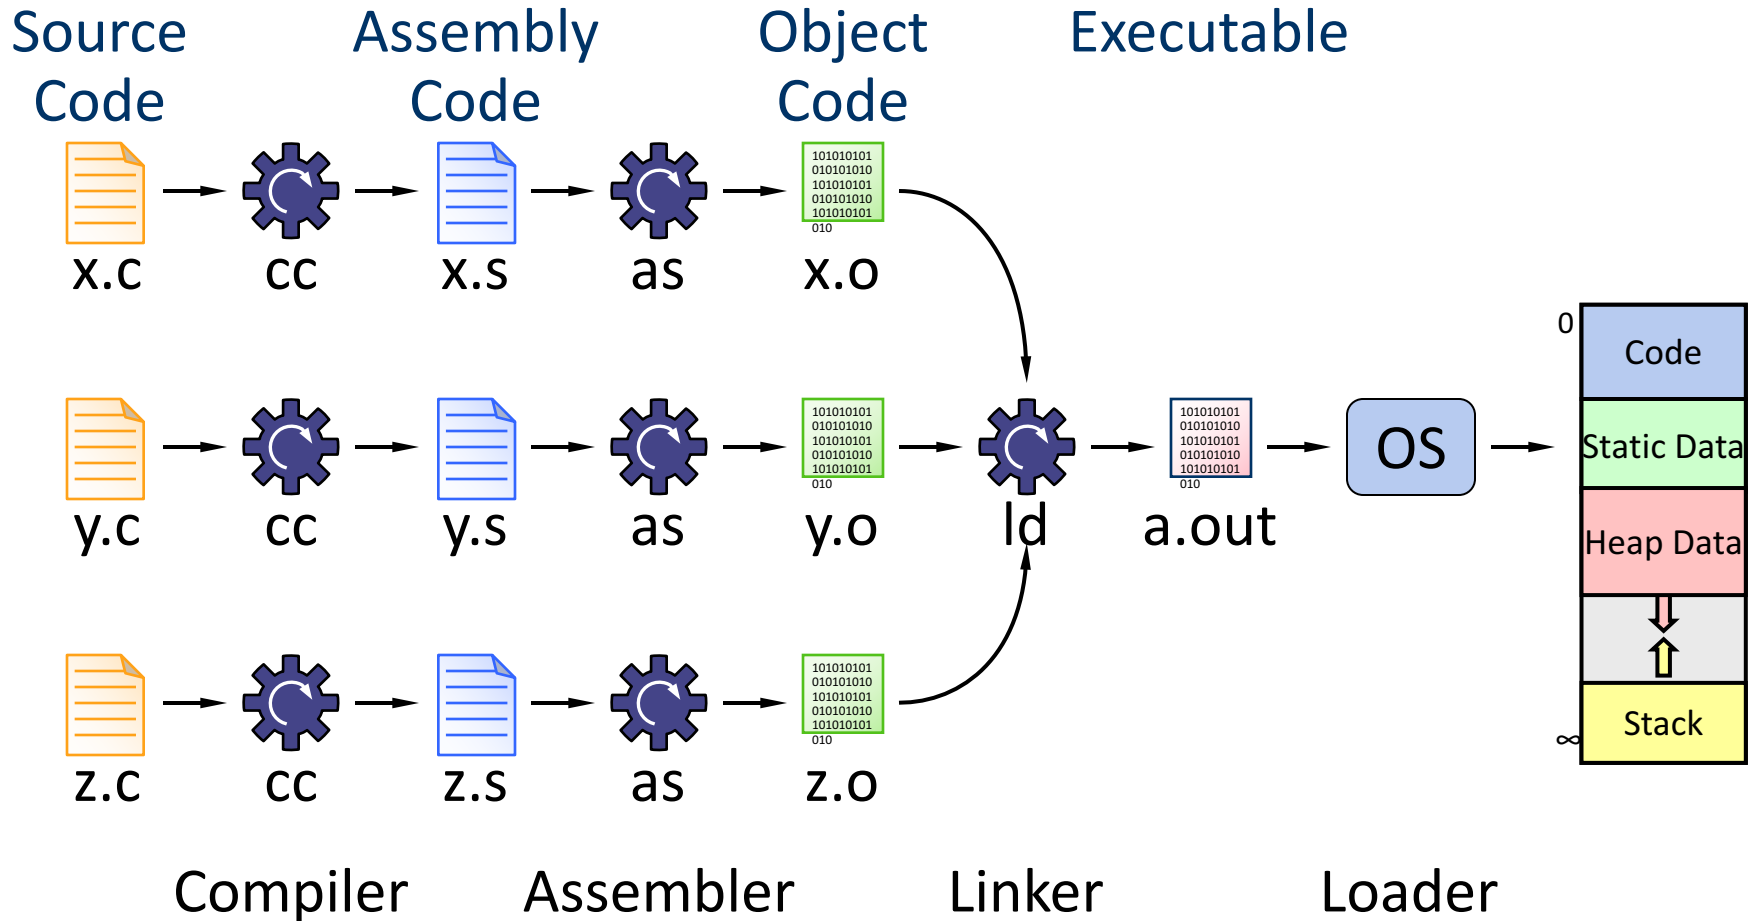
## ❑ How to support functions in assembly

- Passing arguments and getting a return value
- Using the stack and stack frame.
- How to ensure that live values in registers are preserved after a function call.

## ❑ Introduction to linkers and loaders

- Basic relationship of compiler, assembler, linker and loader.
- Object files
  - Symbol tables and relocation tables

# Source to Process Translation



# Linux (ELF—executable and linkable format) object file format

---

*Object files contain more than just machine code instructions!*

**Header:** (of an object file) contains sizes of other parts

**Text:** machine code

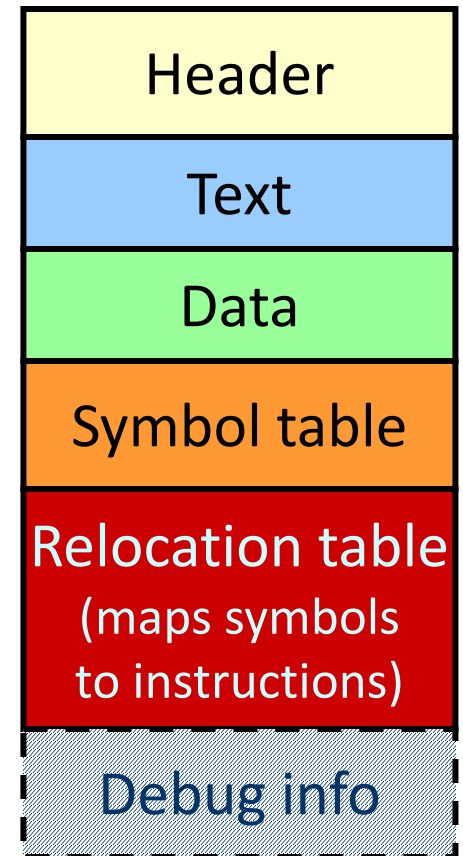
**Data:** global and static data

**Symbol table:** symbols and values

**Relocation table:** references to addresses that may change

**Debug info:** mapping of object back to source (only exists when debugging options are turned on)

## Object code format



# Assembly → Object file - example

## Snippet of C

```
int X = 3;
main() {
    Y = X + 1;
    B();
    ...
}
```

## Snippet of assembly code

```
LDUR    X1, [X27, #0]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
	...		
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [X27, #0] //X27 global reg	
	4	ADDI X9, X1, #1 //X9 local variable Y	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	Location
	X	0	Data
	B	-	-
	main	0	Text
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	X
	12	BL	B

# Class Problem 1

- In the following files, which symbols will be put in the object file's symbol table? If its location is defined in this file, indicate if it would be in the data or text section.

Recall local variables are not in tables:  
b in file 1 and  
\*e in file 2

```
file1.c
extern int bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

file 1 – symbol table

sym	loc
a	data
foo	text
c	-
bar	-

```
file2.c
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

file 2 – symbol table

sym	loc
c	data
bar	text
a	-

# Class Problem 2

Which lines / instructions are in the relocation table for each file?

## file1.c

```
1 extern void bar(int);
2 extern char c[];
3 int a;
4 int foo (int x) {
5     int b;
6     a = c[3] + 1;
7     bar(x);
8     b = 27;
9 }
```

## file 1 - relocation table

line	type	dep
6	ldur	c
6	stur	a
7	bl	bar

## file2.c

```
1 extern int a;
2 char c[100];
3 void bar (int y) {
4     char e[100];
5     a = y;
6     c[20] = e[7];
7 }
```

## file 2 - relocation table

line	type	dep
5	stur	a
6	stur	c

Note: in a real relocation table, the “line” would really be the address in “text” section of the instruction we need to update.



# Some additional questions

---

## file1.c

```
extern void bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

## file2.c

```
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

A) What if file2.c contains `extern int j`, but no reference to `j`?

A smart compiler would not put `j` in symbol table. A dumb one would. It's a benign issue, no harm done if `j` is in symbol table uselessly.

B) What if variable 'e' is static?

It is in the data section (no longer stack) and any reference to it is in relocation table. It is also in the symbol table so the address can be calculated, but it will be flagged not to be exported to other files, since the scope is local.

C) What if the externs in file1.c are deleted?

You should get a compile error

# Loader

---

- ❑ Executable file is sitting on the disk
- ❑ Puts the executable file code image into memory and asks the operating system to schedule it as a new process
  - Creates new address space for program large enough to hold text and data segments, along with a stack segment
  - Copies instructions and data from executable file into the new address space (this may be anywhere in memory)
  - Initializes registers (PC and SP most important)
- ❑ Linking used to be straight forward, but times are changing; it is not simple anymore.
  - We now delay some of the linking to load time
  - Some systems even delay some code optimization (usually a compiler job) to load time
  - Loaders must deal with more sophisticated operating systems

# Things to remember

---

- ❑ Compiler converts a single source code file into a single assembly language file
- ❑ Assembler handles directives (.fill), converts what it can to machine language, and creates a checklist for the linker (relocation table). This changes each .s file into a .o file
- ❑ Assembler does 2 passes to resolve addresses, handling internal forward references
- ❑ Linker combines several .o files and resolves absolute addresses
- ❑ Linker enables separate compilation: Thus unchanged files, including libraries need not be recompiled.
- ❑ Linker resolves remaining addresses.
- ❑ Loader loads executable into memory and begins execution

---

# Floating Point Arithmetic

# Why floating point

---

- ❑ Have to represent real numbers somehow
- ❑ Rational numbers
  - Ok, but can be cumbersome to work with
  - Falls apart for  $\sqrt{2}$  and other irrational numbers
- ❑ Fixed point
  - Do everything in thousandths (or millions, etc.)
  - Not always easy to pick the right units
  - Different scaling factors for different stages of computation
- ❑ Scientific notation: this is good!
  - Exponential notation allows HUGE dynamic range
  - Constant (approximately) relative precision across the whole range

# Floating point before IEEE-754 standard

---

## ❑ Late 1970s formats

- About two dozen different, incompatible floating point number formats
- Precisions from about 4 to about 17 decimal digits
- Ranges from about  $10^{19}$  to  $10^{322}$

## ❑ Sloppy arithmetic

- Last few bits were often wrong, and in different ways
- Overflow sometimes detected, sometimes ignored
- Arbitrary, almost random rounding modes
  - Truncate, round up, round to nearest
- Addition and multiplication not necessarily commutative
  - Small differences due to roundoff errors


# IEEE floating point

---

- ❑ Standard set by IEEE
  - Intel took the lead in 1976 for a good standard
  - First working implementation: Intel 8087 floating point coprocessor, 1980
  - Full formal adoption: 1985
  - Updated in 2008
- ❑ Rigorous specification for high accuracy computation
  - Made every bit count
  - Dependable accuracy even in the lowest bits
  - Predictable, reasonable behavior for exceptional conditions
    - (divide by zero, overflow, etc.)

# IEEE Floating point format (single precision)

- ❑ Sign bit: (0 is positive, 1 is negative)
- ❑ Significand: (also called the *mantissa*; stores the 23 most significant bits after the decimal point)
- ❑ Exponent: used biased base 127 encoding
  - Add 127 to the value of the exponent to encode:
  - $-127 \rightarrow 00000000$        $1 \rightarrow 10000000$
  - $-126 \rightarrow 00000001$        $2 \rightarrow 10000001$
  - ...      ...
  - $0 \rightarrow 01111111$        $128 \rightarrow 11111111$
- ❑ How do you represent zero ? Special convention:
  - Exponent: -127 (all zeroes ), Significand 0 (all zeroes), Sign + or -






# Floating Point Representation

---

$$10.625_{10} \longrightarrow 1010.101_2$$

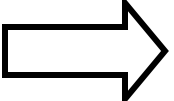
$$1.010101 \times 2^3$$


**This must be a 1!  
So don't store it.**



# Floating Point Representation

---

$10.625_{10}$    $1010.101_2$



**1.010101**  $\times 2^3$

**This must be a 1!  
So don't store it.**

**+/-**

1 bit

**Exponent (3)**

8 bits

**Significand (1010101)**

23 bits

$10.625_{10} = 0 \ 10000010 \ 010101000000000000000000$

# Class Problem

---

- ❑ What is the value (in decimal) of the following IEEE 754 floating point encoded number?

1	10000101	010110010000000000000000
---	----------	--------------------------

# Floating point multiplication

---

- ❑ Add exponents (don't forget to account for the bias of 127)
- ❑ Multiply significands (don't forget the implicit **1** bits)
- ❑ Renormalize if necessary
- ❑ Compute sign bit (simple exclusive-or)

# Floating point multiply

$$10.625_{10} = 1010.101_2 \Rightarrow$$

$$10_{10} = 1010_2 \Rightarrow$$

0	10000010	010101000000000000000000
	+	×
0	10000010	010000000000000000000000
	-127	
0	10000101	101010010000000000000000

	<b>1</b> 0 1 0 1 0 1
×	1 0 1
	1 0 1 0 1 0 1
	1 0 1 0 1 0 1 <b>0 0</b>
	<b>1</b> 1 0 1 0 1 0 0 1

$$1101010.01_2 = 106.25_{10}$$

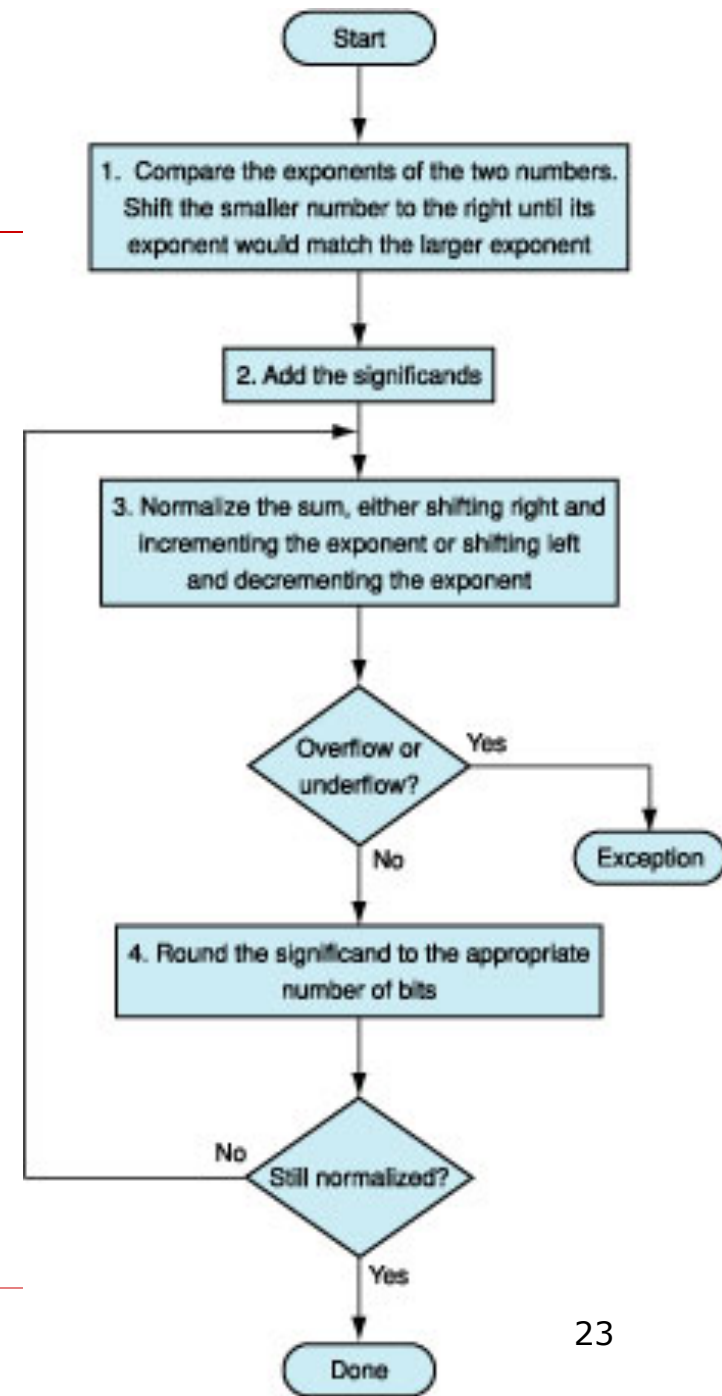
# Floating point addition

---

- ❑ More complicated than floating point multiplication!
- ❑ If exponents are unequal, must shift the significand of the smaller number to the right to align the corresponding place values
- ❑ Once numbers are aligned, simple addition (could be subtraction, if one of the numbers is negative)
- ❑ Renormalize (which could be messy if the numbers had opposite signs; for example, consider addition of +1.5000 and  $-1.4999$ )
- ❑ Added complication: rounding to the correct number of bits to store could denormalize the number, and require one more step

# Floating point Addition

1. Shift smaller exponent right to match larger.
2. Add significands
3. Normalize and update exponent
4. Check for “out of range”



# Class Problem

---

**Show how to add the following 2 numbers using IEEE floating point addition:  $101.125 + 13.75$**



# Class Problem

101.125



13.75



Shift by  $6-3 = 3$

Shift mantissa by difference in exponent



Sum Significands


**1** 1 0 0 1 0 1 0 0 1

+ 0 0 0 **1** 1 0 1 1 1 0

**1** 1 1 0 0 1 0 1 1 1

Sum didn't overflow, so  
no re-normalization  
needed

Note: When shifting to the right, the first  
shift should put the implicit **1**, then 0's



= 114.875

# Class Problem

---

**Show how to add the following 2 numbers using IEEE floating point addition:  $117.125 + 13.75$**

# Class Problem

117.125



13.75



Shift by  $6-3 = 3$


Shift mantissa by difference in exponent



Sum Significands

$$\begin{array}{r}
 1110101001 \\
 + 0001101110 \\
 \hline
 10000010111
 \end{array}$$

Note: When shifting to the right, the first shift should put the implicit **1**, then 0's



= 130.875

Sum overflows, re-normalize by adding one to exponent and shifting mantissa by one

## More precision and range

---

- ❑ We have described IEEE-754 binary32 floating point format, commonly known as “single precision” (“float” in C/C++)
  - 24 bits precision; equivalent to about 7 decimal digits
  - $3.4 * 10^{38}$  maximum value
  - Good enough for most but not all calculations
- ❑ IEEE-754 also defines a larger binary64 format, “double precision” (“double” in C/C++)
  - 53 bits precision, equivalent to about 16 decimal digits
  - $1.8 * 10^{308}$  maximum value
  - Most accurate physical values currently known only to about 47 bits precision, about 14 decimal digits

# Extreme floating point



---

- ❑ binary128, “quad precision”; recent addition to standard:
  - 113 bits precision, about 34 decimal digits
  - $1.2 \times 10^{4932}$  maximum value
  - Very rarely used, but some computations require extreme accuracy to limit cumulative roundoff error
- ❑ Another recent addition was binary16, “half precision”
  - 11 bits precision, about 3.3 decimal digits
  - 65504 maximum value
  - Used in graphics processors for accurate rendering of scenes with a large dynamic range in lighting levels.
  - Minimizes storage per pixel

# Single (“float”) percision

---

## Single Precision



## Double Precision